

```
1 #importing all the important libraries
2 import sklearn.preprocessing
3 import pandas as pd
4 import numpy as np
5 import seaborn as sns
6 import matplotlib.pyplot as plt
7 import warnings
8
9 import sklearn as sk
10 from sklearn import metrics
11 from sklearn.model_selection import train_test_split, ShuffleSplit, learning_cu
12 from sklearn.preprocessing import StandardScaler
13 from sklearn.feature_selection import SelectKBest
14 from sklearn.pipeline import Pipeline
15 from sklearn.metrics import classification_report, confusion_matrix
16 from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
17 from sklearn.linear_model import LogisticRegression
18 from sklearn.svm import SVC
19
20 warnings.filterwarnings('ignore')
```

```
1 df_train=pd.read_csv('/content/drive/MyDrive/ML/ML_A2/Train - Train.csv')
2 df_test=pd.read_csv('/content/drive/MyDrive/ML/ML_A2/test - test.csv')
3 display(df_train.shape)
4 display(df_test.shape)
```

```
(44550, 41)
(14850, 41)
```

```
1 #checking the sample of train data
2
3 df_train.head()
```

```

    id amount_tsh date_recorded funder gps_height installer longitude
1 #unique and null values
2 for col in df_train.columns.values:
3     list_vals = pd.unique(df_train[col]) #list of unique values
4     print('\033[1m' + col + '\033[0m' + ' has ' + str(len(list_vals)) + ' unique '
5     if len(list_vals) < 10:
6         list_str = ''
7         for n in range(0, len(list_vals)):
8             list_str = list_str + str(list_vals[n]) + ', '
9     print('\033[1m' + ' ##### These are: ' + '\033[0m' + list_str[0:len(list_s

```

```

id has 44550 unique values, 0 null entries and datatype int64
amount_tsh has 85 unique values, 0 null entries and datatype float64
date_recorded has 346 unique values, 0 null entries and datatype object
funder has 1652 unique values, 2793 null entries and datatype object
gps_height has 2396 unique values, 0 null entries and datatype int64
installer has 1855 unique values, 2807 null entries and datatype object
longitude has 43155 unique values, 0 null entries and datatype float64
latitude has 43155 unique values, 0 null entries and datatype float64
wpt_name has 28991 unique values, 0 null entries and datatype object
num_private has 59 unique values, 0 null entries and datatype int64
basin has 9 unique values, 0 null entries and datatype object
##### These are: Pangani, Lake Nyasa, Rufiji, Lake Tanganyika, Lake Vic
subvillage has 16618 unique values, 287 null entries and datatype object
region has 21 unique values, 0 null entries and datatype object
region_code has 27 unique values, 0 null entries and datatype int64
district_code has 20 unique values, 0 null entries and datatype int64
lga has 125 unique values, 0 null entries and datatype object
ward has 2080 unique values, 0 null entries and datatype object
population has 956 unique values, 0 null entries and datatype int64
public_meeting has 3 unique values, 2491 null entries and datatype object
##### These are: True, False, nan
recorded_by has 1 unique values, 0 null entries and datatype object
##### These are: GeoData Consultants Ltd
scheme_management has 13 unique values, 2832 null entries and datatype object
scheme_name has 2507 unique values, 21110 null entries and datatype object
permit has 3 unique values, 2336 null entries and datatype object
##### These are: True, False, nan
construction_year has 55 unique values, 0 null entries and datatype int64
extraction_type has 18 unique values, 0 null entries and datatype object
extraction_type_group has 13 unique values, 0 null entries and datatype object
extraction_type_class has 7 unique values, 0 null entries and datatype object
##### These are: gravity, handpump, motorpump, submersible, other, rope
management has 12 unique values, 0 null entries and datatype object
management_group has 5 unique values, 0 null entries and datatype object
##### These are: user-group, commercial, parastatal, unknown, other
payment has 7 unique values, 0 null entries and datatype object
##### These are: pay per bucket, never pay, pay annually, pay monthly,
payment_type has 7 unique values, 0 null entries and datatype object
##### These are: per bucket, never pay, annually, monthly, unknown, on
water_quality has 8 unique values, 0 null entries and datatype object
##### These are: soft, salty abandoned, unknown, salty, fluoride, milky
quality_group has 6 unique values, 0 null entries and datatype object
##### These are: good, salty, unknown, fluoride, milky, colored
quantity has 5 unique values, 0 null entries and datatype object
##### These are: enough, insufficient, dry, unknown, seasonal
quantity_group has 5 unique values, 0 null entries and datatype object
##### These are: enough, insufficient, dry, unknown, seasonal

```

```

##### These are: enough, insufficient, dry, unknown, seasonal
source has 10 unique values, 0 null entries and datatype object
source_type has 7 unique values, 0 null entries and datatype object
##### These are: spring, river/lake, shallow well, borehole, rainwater
source_class has 3 unique values, 0 null entries and datatype object
##### These are: groundwater, surface, unknown
waterpoint_type has 7 unique values, 0 null entries and datatype object
##### These are: communal standpipe, hand pump, other, communal standpipe
waterpoint_type_group has 6 unique values, 0 null entries and datatype object
##### These are: communal standpipe, hand pump, other, cattle trough, in
status_group has 3 unique values, 0 null entries and datatype object

```

▼ DATA PREPROCESSING

```

1 #Dropping of the variable which has too many numerical value
2 #since it will not play significant role in the output
3
4
5 #dropping the variable wpt_name since it has too many distinct numerical value
6 #Dropping the wpt_name variable more than 50% different values
7 df_train.drop(['wpt_name'], axis=1, inplace=True)
8 df_test.drop(['wpt_name'], axis=1, inplace=True)

```

```

1 #Dropping the id variable since it does not have any significance in the water
2 df_train.drop(['id'], axis=1, inplace=True)
3 df_test.drop(['id'], axis=1, inplace=True)

```

```

1 #DATA PREPROCESSING
2 #we have date recorded as datatype object converting it to integer
3 #converting the date format dd-mm-yyyy to yyyymmdd and converting it to the integer
4 #for training
5 df_train['date_recorded'] = pd.to_datetime(df_train['date_recorded']).dt.strftime('%Y%m%d')
6 df_train['date_recorded'] = df_train['date_recorded'].astype(int)
7 print( 'datatype od df_test changed to ', df_train['date_recorded'].dtype )
8
9 #for testing
10 df_test['date_recorded'] = pd.to_datetime(df_test['date_recorded']).dt.strftime('%Y%m%d')
11 df_test['date_recorded'] = df_test['date_recorded'].astype(int)
12 print( 'datatype od df_test changed to ',df_test['date_recorded'].dtype )

```

```

datatype od df_test changed to int64
datatype od df_test changed to int64

```

```

1 df_train['date_recorded'].nunique() #no. of unique values in date recorded

346

```

```


1 # Find columns with 'object' dtypes.
2 list(df_train.select_dtypes(np.number))

```

```
['amount_tsh',  
 'date_recorded',  
 'gps_height',  
 'longitude',  
 'latitude',  
 'num_private',  
 'region_code',  
 'district_code',  
 'population',  
 'construction_year']
```

initially checking the correlation between variables to
amongst the numerical values

```
1 # visualization of correlation by heatmap for red wine  
2  
3 corr = df_train.corr(method = 'spearman')  
4 fig, ax = plt.subplots(figsize = (10,10))  
5 sns.heatmap(abs(corr), annot = True)  
6 plt.show()
```



From above its comes to notice that Most of the numerical variable columns are not correlated with each other.



```
1 # Find columns with 'object' dtypes.
2 list(df_train.select_dtypes(exclude=[np.number]))
```

```
['funder',
 'installer',
 'basin',
 'subvillage',
 'region',
 'lga',
 'ward',
 'public_meeting',
 'recorded_by',
 'scheme_management',
 'scheme_name',
 'permit',
 'extraction_type',
 'extraction_type_group',
 'extraction_type_class',
 'management',
 'management_group',
 'payment',
 'payment_type',
 'water_quality',
 'quality_group',
 'quantity',
 'quantity_group',
 'source',
 'source_type',
 'source_class',
 'waterpoint_type',
 'waterpoint_type_group',
 'status_group']
```

features having **Non-Numeric Value** as content are:

```
['funder','installer','basin','subvillage','region','lga',
 'ward','public_meeting','recorded_by','scheme_management','scheme_name','permit','extraction_type',
 'extraction_type_group','extraction_type_class','management', 'management_group',
 'payment','payment_type','water_quality','quality_group', 'quantity','quantity_group', 'source',
 'source_type', 'source_class','waterpoint_type','waterpoint_type_group', 'status_group']
```

▼ Missing value

funder has , 2793 null entries and datatype object
 installer , 2807 null entries and datatype object
 subvillage, 287 null entries and datatype object
 public_meeting , 2491 null entries and datatype object
 scheme_management , 2832 null entries and datatype object
 scheme_name , 21110 null entries and datatype object
 permit has , 2336 null entries and datatype object

Imputation of the missing values with most frequent values using **mode** value of the column.

And cross checking after the imputation

```
1 #Imputation of the missing values with most frequent values using mode.
2 #Imputing the missing values
3 df_train['permit'].fillna(df_train['permit'].mode()[0], inplace=True)
4 df_train['funder'].fillna(df_train['funder'].mode()[0], inplace=True)
5 df_train['scheme_management'].fillna(df_train['scheme_management'].mode()[0], in
6 df_train['public_meeting'].fillna(df_train['public_meeting'].mode()[0], inplace=
7 df_train['subvillage'].fillna(df_train['subvillage'].mode()[0], inplace=True)
8 df_train['installer'].fillna(df_train['installer'].mode()[0], inplace=True)
9 df_train['scheme_name'].fillna(df_train['scheme_name'].mode()[0], inplace=True)
```

```
1 #Imputation of the missing values with most frequent values using mode.
2 #Imputing the missing values
3 df_test['permit'].fillna(df_test['permit'].mode()[0], inplace=True)
4 df_test['funder'].fillna(df_test['funder'].mode()[0], inplace=True)
5 df_test['scheme_management'].fillna(df_test['scheme_management'].mode()[0], inp
6 df_test['public_meeting'].fillna(df_test['public_meeting'].mode()[0], inplace=T
7 df_test['subvillage'].fillna(df_test['subvillage'].mode()[0], inplace=True)
8 df_test['installer'].fillna(df_test['installer'].mode()[0], inplace=True)
9 df_test['scheme_name'].fillna(df_test['scheme_name'].mode()[0], inplace=True)
```

```
1 #Dropping of the unnecessary column from the test data
2 df_test.drop(['Unnamed: 0'], axis=1, inplace=True)
```

```
1 #After the imputation of the vlues in the columns of the dataframe checking if
2 df_test.isnull().sum()
```

```
amount_tsh      0
date_recorded   0
funder           0
gps_height      0
installer       0
longitude       0
latitude        0
num_private     0
```

```
basin 0
subvillage 0
region 0
region_code 0
district_code 0
lga 0
ward 0
population 0
public_meeting 0
recorded_by 0
scheme_management 0
scheme_name 0
permit 0
construction_year 0
extraction_type 0
extraction_type_group 0
extraction_type_class 0
management 0
management_group 0
payment 0
payment_type 0
water_quality 0
quality_group 0
quantity 0
quantity_group 0
source 0
source_type 0
source_class 0
waterpoint_type 0
waterpoint_type_group 0
dtype: int64
```

```
1 df_train.shape
```

```
(44550, 39)
```

```
1 #encoding with ordinal encoder to train data
2 from sklearn.preprocessing import OrdinalEncoder
3 enc = OrdinalEncoder()
4 for col in df_train.columns:
5     df_train[col] = enc.fit_transform(df_train[[col]])
```


```
1 #for test data
2 for col in df_test.columns:
3     df_test[col] = enc.fit_transform(df_test[[col]])
```

```
1 #Dropping the variables
2 df_test.drop(['recorded_by'], axis=1, inplace=True)
3
```

```
1 df_train.drop(['recorded_by'], axis=1, inplace=True)
```

```
1 #Correlation of the newly encoded variables
```

```
2 # visualization of correlation by heatmap
3
4 corr = df_train.corr(method = 'spearman')
5 fig, ax = plt.subplots(figsize = (20,20))
6 sns.heatmap(abs(corr), annot = False)
7 plt.show()
```

```
1 #extraction type
2 df_train.columns.values
```

```
array(['amount_tsh', 'date_recorded', 'funder', 'gps_height', 'installer',
      'longitude', 'latitude', 'num_private', 'basin', 'subvillage',
      'region', 'region_code', 'district_code', 'lga', 'ward',
      'population', 'public_meeting', 'scheme_management', 'scheme_name',
      'permit', 'construction_year', 'extraction_type',
      'extraction_type_group', 'extraction_type_class', 'management',
      'management_group', 'payment', 'payment_type', 'water_quality',
      'quality_group', 'quantity', 'quantity_group', 'source',
      'source_type', 'source_class', 'waterpoint_type',
      'waterpoint_type_group', 'status_group'], dtype=object)
```



```
1 df_test.columns.values
```

```
array(['amount_tsh', 'date_recorded', 'funder', 'gps_height', 'installer',
      'longitude', 'latitude', 'num_private', 'basin', 'subvillage',
      'region', 'region_code', 'district_code', 'lga', 'ward',
      'population', 'public_meeting', 'scheme_management', 'scheme_name',
      'permit', 'construction_year', 'extraction_type',
      'extraction_type_group', 'extraction_type_class', 'management',
      'management_group', 'payment', 'payment_type', 'water_quality',
      'quality_group', 'quantity', 'quantity_group', 'source',
      'source_type', 'source_class', 'waterpoint_type',
      'waterpoint_type_group'], dtype=object)
```

```
1 #Dropping the variables extraction_type_class , and group has near to unity cor
2 df_train.drop(['extraction_type_group'], axis=1, inplace=True)
3 df_train.drop(['extraction_type_class'], axis=1, inplace=True)
```

```
1 #test Dropping the variables extraction_type_class , and group has near to unity
2 df_test.drop(['extraction_type_group'], axis=1, inplace=True)
3 df_test.drop(['extraction_type_class'], axis=1, inplace=True)
```

```

1 #train test Dropping the variables quantity_group has near to unity correlation
2 df_train.drop(['quantity_group'], axis=1, inplace=True)
3 df_test.drop(['quantity_group'], axis=1, inplace=True)

1 #train test Dropping the variables source_type and source has near to unity correlation
2 df_train.drop(['source_type'], axis=1, inplace=True)
3 df_test.drop(['source_type'], axis=1, inplace=True)

1 #train test Dropping the variables waterpoint_type_group and waterpoint_type has near to unity correlation
2 df_train.drop(['waterpoint_type_group'], axis=1, inplace=True)
3 df_test.drop(['waterpoint_type_group'], axis=1, inplace=True)

1 display(df_train.shape)
2 display(df_test.shape)

(44550, 33)
(14850, 32)

1 #train data
2 X = pd.DataFrame(df_train.iloc[:,0:-1])
3 Y = pd.DataFrame(df_train['status_group'])

1 #test data processed
2 Xt = pd.DataFrame(df_test)

1 from sklearn.feature_selection import RFECV
2 from sklearn.model_selection import GridSearchCV
3 from sklearn.model_selection import train_test_split
4 from sklearn.ensemble import RandomForestClassifier
5
6 X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.33, random_state=42)
7

```

▼ 1.0 RFClassification

```

1 #RFC
2 #Hyperparameter tuning
3 rf = RandomForestClassifier(criterion='gini', n_estimators=500, max_features='auto')
4 param_grid = {"min_samples_split" : [4, 6, 8], "n_estimators" : [500, 700, 1000]}
5
6 gs = GridSearchCV(estimator=rf, param_grid=param_grid, scoring='accuracy', cv=2, n_jobs=-1)
7
8 gs = gs.fit(X_train, y_train.values.ravel())
9
10 print(gs.best_score_)
11 print(gs.best_params_)

0.7888970785312248

```

```
{'min_samples_split': 6, 'n_estimators': 700}
```

```
1 gs.best_estimator_.feature_importances_
```

```
array([0.02624758, 0.0505222 , 0.0342603 , 0.04281361, 0.02975738,
        0.07936963, 0.07577333, 0.00122261, 0.01328392, 0.05208837,
        0.01676589, 0.0150217 , 0.01824166, 0.02666965, 0.03985745,
        0.03146935, 0.00521997, 0.01155649, 0.02730021, 0.00587742,
        0.04370495, 0.04067509, 0.01491027, 0.0058288 , 0.01897732,
        0.01638281, 0.01035214, 0.01162825, 0.1380757 , 0.02323805,
        0.00654654, 0.06636136])
```

```
1 rf32 = RandomForestClassifier(criterion='gini',min_samples_split=6,n_estimators=
```

```
2
```

```
3 rf32.fit(X_train, y_train.values.ravel())
```

```
4 print('RFC has oob score after 1st elimination is :',rf32.oob_score_)
```

```
5
```

```
RFC has oob score after 1st elimination is : 0.8056486196730099
```

```
1 prediction = rf32.predict(X_test)
```

```
2 print('Accuracy for RFC is', metrics.accuracy_score(y_test, prediction))
```

```
Accuracy for RFC is 0.8000952251394368
```

```
1 rf32.feature_importances_
```

```
array([0.02624758, 0.0505222 , 0.0342603 , 0.04281361, 0.02975738,
        0.07936963, 0.07577333, 0.00122261, 0.01328392, 0.05208837,
        0.01676589, 0.0150217 , 0.01824166, 0.02666965, 0.03985745,
        0.03146935, 0.00521997, 0.01155649, 0.02730021, 0.00587742,
        0.04370495, 0.04067509, 0.01491027, 0.0058288 , 0.01897732,
        0.01638281, 0.01035214, 0.01162825, 0.1380757 , 0.02323805,
        0.00654654, 0.06636136])
```

▼ 2.0 XGBoost Classification

```
1 from xgboost import XGBClassifier
```

```
2 from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
```

```
3 from sklearn.model_selection import StratifiedKFold
```

```
4 xgb = XGBClassifier(learning_rate=0.02, n_estimators=600, objective='binary:log
```

```
5 folds = 3
```

```
6 param_comb = 5
```

```
7 skf = StratifiedKFold(n_splits=folds, shuffle = True, random_state = 1001)
```

```
8 scorer = sklearn.metrics.make_scorer(sklearn.metrics.f1_score, average = 'weigh
```

```
9 params = {'min_child_weight': [1, 5, 10], 'gamma': [0.5, 1, 1.5, 2, 5], 'subsample
```

```
10 random_search = RandomizedSearchCV(xgb, param_distributions=params, n_iter=param
```

```
11 random_search.fit(X_train, y_train )
```

```
Fitting 3 folds for each of 5 candidates, totalling 15 fits
```

```
[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=4)]: Done 15 out of 15 | elapsed: 10.5min finished
RandomizedSearchCV(cv=<generator object _BaseKFold.split at 0x7fe66e890c50>,
                  error_score=nan,
                  estimator=XGBClassifier(base_score=0.5, booster='gbtree',
                                          colsample_bylevel=1,
                                          colsample_bynode=1,
                                          colsample_bytree=1, gamma=0,
                                          learning_rate=0.02,
                                          max_delta_step=0,
                                          max_depth=3, min_child_weight=1,
                                          missing=None, n_estimators=600,
                                          n_jobs=1, nthread=1,
                                          objective='binary:logist...
                                          seed=None, silent=True,
                                          subsample=1,
                                          verbosity=1),
                  iid='deprecated', n_iter=5, n_jobs=4,
                  param_distributions={'colsample_bytree': [0.6, 0.8, 1.0],
                                      'gamma': [0.5, 1, 1.5, 2, 5],
                                      'max_depth': [3, 4, 5],
                                      'min_child_weight': [1, 5, 10],
                                      'subsample': [0.6, 0.8, 1.0]},
                  pre_dispatch='2*n_jobs', random_state=1001, refit=True,
                  return_train_score=False,
                  scoring=make_scorer(f1_score, average=weighted),
                  verbose=3)
```

```
1 print('best score xgb ', random_search.best_score_)
2 print('xgb best parameters ', random_search.best_params_)
3 print('best indexx xgb ', random_search.best_index_)
4 print('best estimator xgb ', random_search.best_estimator_)
```

```
best score xgb 0.75399236221129
xgb best parameters {'subsample': 0.6, 'min_child_weight': 1, 'max_depth': 5}
best indexx xgb 1
best estimator xgb XGBClassifier(base_score=0.5, booster='gbtree', colsample
                               colsample_bynode=1, colsample_bytree=0.8, gamma=1.5,
                               learning_rate=0.02, max_delta_step=0, max_depth=5,
                               min_child_weight=1, missing=None, n_estimators=600, n_jobs=1,
                               nthread=1, objective='multi:softprob', random_state=0,
                               reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                               silent=True, subsample=0.6, verbosity=1)
```

```
1 #Applying the best parameters of the xgboost to model and testing the model on
2 xgb32=random_search.best_estimator_
3
4 xgb32.fit(X_train, y_train.values.ravel())
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=0.8, gamma=1.5,
              learning_rate=0.02, max_delta_step=0, max_depth=5,
              min_child_weight=1, missing=None, n_estimators=600, n_jobs=1,
              nthread=1, objective='multi:softprob', random_state=0,
              reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
              silent=True, subsample=0.6, verbosity=1)
```

```
1 xgb32.feature_importances_
```

```
array([0.0462949 , 0.01619986, 0.01752931, 0.01707107, 0.01616328,
       0.02376452, 0.01640674, 0.01277481, 0.02424673, 0.01005085,
       0.02871218, 0.02235642, 0.01806406, 0.02134535, 0.01183366,
       0.0141106 , 0.02705525, 0.01536601, 0.01416728, 0.0191035 ,
       0.02868589, 0.03501341, 0.03257824, 0.02063816, 0.03585385,
       0.02510552, 0.01965024, 0.02040447, 0.20942442, 0.03958386,
       0.02705507, 0.11339048], dtype=float32)
```

```
1 prediction = xgb32.predict(X_test)
```

```
2 print('Accuracy xgb is', metrics.accuracy_score(y_test, prediction))
```

```
Accuracy xgb is 0.7749285811454224
```

▼ 3.0 KNN

```
1
2 from sklearn.neighbors import KNeighborsClassifier
3 from sklearn.pipeline import Pipeline
4 k_range = list(range(5,10))
5
6 param_grid = dict(n_neighbors=k_range)
7
8 pipe = Pipeline([('sc', StandardScaler()), ('knn', KNeighborsClassifier(algorithm='brute',
9 params = {'knn__n_neighbors': k_range }
10 clf = GridSearchCV(estimator=pipe, param_grid=params, cv=5, return_train_score=True)
11 clf.fit(X_train, y_train)
```

```
GridSearchCV(cv=5, error_score=nans,
              estimator=Pipeline(memory=None,
                                 steps=[('sc',
                                          StandardScaler(copy=True,
                                                             with_mean=True,
                                                             with_std=True)),
                                          ('knn',
                                           KNeighborsClassifier(algorithm='brute',
                                                                  leaf_size=30,
                                                                  metric='minkowski',
                                                                  metric_params=None,
                                                                  n_jobs=None,
                                                                  n_neighbors=5,
                                                                  p=2,
                                                                  weights='uniform'))]),
              verbose=False),
              iid='deprecated', n_jobs=None,
              param_grid={'knn__n_neighbors': [5, 6, 7, 8, 9]},
```

```

pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring=None, verbose=0)

1 print('best score knn ',clf.best_score_ )
2 print('knn best parameters ',clf.best_params_ )
3 print('best index knn ', clf.best_index_ )
4 print('best estimator knn ',clf.best_estimator_ )

best score knn 0.7319751603272407
knn best parameters {'knn__n_neighbors': 5}
best index knn 0
best estimator knn Pipeline(memory=None,
                             steps=[('sc',
                                       StandardScaler(copy=True, with_mean=True, with_std=True)),
                                      ('knn',
                                       KNeighborsClassifier(algorithm='brute', leaf_size=30,
                                                            metric='minkowski', metric_params=None,
                                                            n_jobs=None, n_neighbors=5, p=2,
                                                            weights='uniform'))],
                             verbose=False)

```

Observations We are having 3 frameworks for prediction of the

1. RandomForestClassifier having accuracy of **80** percent approx
2. XGBoost Classifier having the accuracy of the **78** percent approx
3. KNN Classifier having Accuracy of **77** percent

So we will select the RFCClassification for the further round of the best feature extraction

1.1 RFC Round 2 of feature elimination.

2. Checking the importance of a particular feature with respect to the feature importance score

```

1 #Round 2 of feature elimination
2 #Checking the importance of the features with respect to the feature importance
3 RFC_feature=rf32.feature_importances_
4 column_name=X.columns
5 dff=pd.DataFrame(X.columns,rf32.feature_importances_)
6 val=pd.DataFrame(RFC_feature)
7 nam=pd.DataFrame(column_name)
8 nam.rename( columns={0 : 'Feaures'}, inplace=True )

```

```
9 val.rename( columns={0 : 'Importance'}, inplace=True )  
10 df1 = pd.concat([val,nam], axis=1)  
11 df1.sort_values(by=[ 'Importance'])
```

	Importance	Feaures
7	0.001223	num_private
16	0.005220	public_meeting

1.1.1 Round 2 Feature elimination

As from above it is clear that the minimum importance is of the variable

1. 0.001223 num_private
2. 0.005220 public_meeting
3. 0.005829 management_group
4. 0.005877 permit
5. 0.006547 source_class

-- 0.001223 num_private

```
1 #storing the original dataset to new variable for feature elimination
2 X2=X
3 Y2=Y
4 Xt2=Xt
```

12 0.018242 district_code

```
1 #31
2 X2.drop(['num_private'], axis=1, inplace=True)
3 Xt2.drop(['num_private'], axis=1, inplace=True)
```

```
1 #30
2 X2.drop(['public_meeting'], axis=1, inplace=True)
3 Xt2.drop(['public_meeting'], axis=1, inplace=True)
```

```
1 #29
2 X2.drop(['management_group'], axis=1, inplace=True)
3 Xt2.drop(['management_group'], axis=1, inplace=True)
```

```
1 #28
2 X2.drop(['permit'], axis=1, inplace=True)
3 Xt2.drop(['permit'], axis=1, inplace=True)
```

9 0.072017 gps_height

```
1 #27
2 X2.drop(['source_class'], axis=1, inplace=True)
3 Xt2.drop(['source_class'], axis=1, inplace=True)
```

9 0.052088 subvillage

```
1 X2.shape
```

(44550, 27)

5 0.079370 longitude

```
1 Xt2.shape
```



```
(14850, 27)
```

```
1 X1_train, X1_test, y1_train, y1_test = train_test_split(X2, Y2, test_size=0.33,

1 #RFC
2 #Hyperparameter tuning
3 rf = RandomForestClassifier(criterion='gini',n_estimators=500,max_features='auto')
4 param_grid = {"min_samples_split" : [4, 6, 8],"n_estimators" : [500, 700, 1000]}
5
6 gs1 = GridSearchCV(estimator=rf,param_grid=param_grid,scoring='accuracy',cv=2,n
7
8 gs1 = gs1.fit(X1_train, y1_train.values.ravel())
9
10 print(gs1.best_score_)
11 print(gs1.best_params_)

0.7888300723666577
{'min_samples_split': 8, 'n_estimators': 500}
```

```
1 rf321 = RandomForestClassifier(criterion='gini',min_samples_split=8,n_estimators=
2
3 rf321.fit(X1_train, y1_train.values.ravel())
4 print('oob score RFC after r2 elimination is',rf321.oob_score_)
```

```
oob score RFC after r2 elimination is 0.8059836504958456
```

```
1 prediction = rf321.predict(X1_test)
2 print('Accuracy is', metrics.accuracy_score(y1_test, prediction))
```

```
Accuracy is 0.7996190994422527
```

```
1 p=rf321.feature_importances_
```

```
1 #Currently after round 2 of feature elimination we have pretty much the
```

Checking the importance of a particular feature with respect to the feature importance score

```
1 #Round 3 checking of feature elimination
2 #Checking the importance of the features with respect to the feature importance
3 RFC_feature=rf321.feature_importances_
4 column_name=X2.columns
5 dff=pd.DataFrame(X.columns,rf321.feature_importances_)
6 val=pd.DataFrame(RFC_feature)
7 nam=pd.DataFrame(column_name)
8 nam.rename( columns={0 : 'Feaures'}, inplace=True )
9 val.rename( columns={0 : 'Importance'}, inplace=True )
```

```
10 df1 = pd.concat([val,nam], axis=1)
11 df1.sort_values(by=[ 'Importance' ])
```

	Importance	Feaures
23	0.010668	quality_group
22	0.010682	water_quality
15	0.012704	scheme_management
7	0.013083	basin
10	0.015136	region_code
21	0.016205	payment_type
19	0.016685	management
9	0.017143	region
11	0.017885	district_code
20	0.018762	payment
25	0.026157	source
0	0.027088	amount_tsh
16	0.027700	scheme_name
12	0.028662	lga
4	0.029721	installer
14	0.030869	population
2	0.034615	funder
13	0.039538	ward
3	0.041569	gps_height
18	0.043368	extraction_type
17	0.044465	construction_year
8	0.049241	subvillage
1	0.050352	date_recorded
26	0.073274	waterpoint_type
6	0.074597	latitude
5	0.078750	longitude
24	0.151078	quantity

After feature reduction 2 accuracy is pretty much same

Accuracy is 0.7996190994422527 ,slightly deceased \ oob score is 0.8059836504958456
 Trvina to reduce further more features.

▼ 1.1.2 Features to eliminate for round 3

0.010350 quality_group

22 0.010504 water_quality

15 0.012171 scheme_management

7 0.012798 basin

10 0.014712 region_code

1 X3=X2

2 Y3=Y2

3 Xt3=Xt2

1 #26

2 X3.drop(['quality_group'], axis=1, inplace=True)

3 Xt3.drop(['quality_group'], axis=1, inplace=True)

4

5 #25

6 X3.drop(['water_quality'], axis=1, inplace=True)

7 Xt3.drop(['water_quality'], axis=1, inplace=True)

8

9 #24

10 X3.drop(['scheme_management'], axis=1, inplace=True)

11 Xt3.drop(['scheme_management'], axis=1, inplace=True)

12

13 #23

14 X3.drop(['basin'], axis=1, inplace=True)

15 Xt3.drop(['basin'], axis=1, inplace=True)

16

17 #22

18 X3.drop(['region_code'], axis=1, inplace=True)

19 Xt3.drop(['region_code'], axis=1, inplace=True)

20

21 #X3.shape is (44550, 22)

22 #Xt3.shape is (14850, 22)

▼ Applying the RFC Framework again by tuning the model
 and cross checking with the data

1 #Round 3 of tuning of parameters

2 X2_train, X2_test, y2_train, y2_test = train_test_split(X3, Y3, test_size=0.33,

3

4 #RFC

5 #Hyperparameter tuning

```

6 rf = RandomForestClassifier(criterion='gini',n_estimators=500,max_features='auto')
7 param_grid = {"min_samples_split" : [4, 6, 8],"n_estimators" : [500, 700, 1000]}
8
9 gs2 = GridSearchCV(estimator=rf,param_grid=param_grid,scoring='accuracy',cv=2,n
10
11 gs2= gs2.fit(X2_train, y2_train.values.ravel())
12
13 print(gs2.best_score_)
14 print(gs2.best_params_)

```

```

0.7896341463414633
{'min_samples_split': 6, 'n_estimators': 1000}

```

```

1 rf3 = RandomForestClassifier(criterion='gini',min_samples_split=6,n_estimators=
2
3 rf3.fit(X2_train, y2_train.values.ravel())
4 print('after 3 round of elimination oob score is : ', rf3.oob_score_)

```

```

after 3 round of elimination oob score is : 0.8053135888501742

```

```

1 prediction = rf3.predict(X2_test)
2 print('Accuracy after 3rd round of elimination is', metrics.accuracy_score(y_te

```

```

Accuracy after 3rd round of elimination is 0.798938919874847

```

Observation After reducing the 5 more features out of 27 remaining feature we are getting pretty much same oob score and the accuracy of the model is increased very slightly

➤ As we got the accuracy highest for the RFC

Accuracy of RFC initially(32 features): 0.8000952251394368 ,oob score is :0.8056486196730099

After 2nd round of elimination(27 features): 0.7996190994422527 ,oob score is :0.8059836504958456

After 3rd round of elimination(22 features) : 0.798938919874847 , oob score is :0.8053135888501742

This shows we should consider 1st model only.

▼ Prediction of the output after round 3 of elimination

We are left with total 22 features almost the accuracy of all the models is same so considering the one

```
1 final_prediction=rf3.predict(Xt3)
2 final_prediction=pd.DataFrame(final_prediction)
3 final_prediction
```

	0
0	2.0
1	2.0
2	0.0
3	2.0
4	0.0
...	...
14845	0.0
14846	0.0
14847	2.0
14848	0.0
14849	0.0

14850 rows × 1 columns

```
1 final_prediction.columns[0]
```

```
1 di = {0: "functional", 1: "functional needs repair", 2: "non functional"}
2 final=final_prediction.replace({final_prediction.columns[0] : di })
3 final.to_csv('final_203079016.csv', header=False)
```

```
1 Y=Y.replace({Y.columns[0] : di })
2 df2 = pd.concat([X,Y], axis=1)
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-1-cfb5bd20ea9c> in <module>()
----> 1 Y=Y.replace({Y.columns[0] : di })
      2 df2 = pd.concat([X,Y], axis=1)
```

NameError: name 'Y' is not defined

SEARCH STACK OVERFLOW

```
1 from sklearn.manifold import TSNE
2 import seaborn as sns
3 m= TSNE(learning_rate=50)
4 tsne_features=m.fit_transform(X)
5 tsne_features[1:4,:]
6 df['x']= tsne_features[:,0]
7 df['x']= tsne_features[:,1]
8
```

